

Homework 1

Rento Saijo

Department of Mathematics, Connecticut College

STA336: Statistical Machine Learning

Yan Zhuang, Ph.D.

January 30th, 2026

Disclosure

ChatGPT-5.2 was used to create the YAML portion to format the text nicely; I looked into the documentation for each of the package it used and added/removed unnecessary formattings. See the original RMD file [here](#).

Problem 1

A very flexible approach has the advantage that it can represent a much wider range of possible shapes for f , and thus capture complicated (often non-linear) relationships between predictors X and a response Y . In contrast, a restrictive method like linear regression can only produce linear functions, e.g., $f(X) = \beta_0 + \sum_{j=1}^p \beta_j X_j$. The drawback of high flexibility is reduced interpretability: the fitted \hat{f} can become so complex that it is difficult to understand how any individual predictor X_j is associated with Y , making flexible methods less attractive when inference and interpretability are the goal. Therefore, more flexible approaches are generally preferred when interpretability is not a priority and prediction is the primary objective, since we are willing to trade a clear description of predictor-response relationships for the ability to fit complex patterns; however, even for prediction, the most flexible model is not always best because highly flexible methods can overfit, so a less flexible method can sometimes yield better test performance. Conversely, a less flexible approach is preferred when inference is the goal because restrictive models are much more interpretable. *Source: ISLR2 §2.1.3, p. 24-6.*

Problem 2 (a)

```
# Load libraries.

suppressMessages(library(tidyverse))
suppressMessages(library(GGally))
suppressMessages(library(ISLR2))

# Load data.

data(Auto)

# Count missing values.

colSums(is.na(Auto)) # It seems that we have no missing values.
```

	mpg	cylinders	displacement	horsepower	weight	acceleration
##	0	0	0	0	0	0
##	year	origin	name			

```

##          0          0          0

# Check structure.

tibble::glimpse(Auto)

## Rows: 392

## Columns: 9

## $ mpg      <dbl> 18, 15, 18, 16, 17, 15, 14, 14, 14, 15, 15, 14, 15, 14, 2~
## $ cylinders <int> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 4, 6, 6, 6, 4, ~
## $ displacement <dbl> 307, 350, 318, 304, 302, 429, 454, 440, 455, 390, 383, 34~
## $ horsepower <int> 130, 165, 150, 150, 140, 198, 220, 215, 225, 190, 170, 16~
## $ weight     <int> 3504, 3693, 3436, 3433, 3449, 4341, 4354, 4312, 4425, 385~
## $ acceleration <dbl> 12.0, 11.5, 11.0, 12.0, 10.5, 10.0, 9.0, 8.5, 10.0, 8.5, ~
## $ year       <int> 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 7~
## $ origin     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 3, 1, 1, 1, 3, ~
## $ name       <fct> chevrolet chevelle malibu, buick skylark 320, plymouth sa~

```

In the Auto data set, `mpg` is a quantitative variable but it is typically the response, not a predictor. The quantitative predictors are therefore `cylinders`, `displacement`, `horsepower`, `weight`, `acceleration`, and `year`. The qualitative predictors are `origin` (a categorical variable encoded numerically) and potentially `name` (more on `name` in part (e)).

Problem 2 (b)

```

# Select quantitative predictors.

Auto_quant_preds <- Auto %>%
  dplyr::select(cylinders, displacement, horsepower, weight, acceleration, year)

# Compute range for each quantitative predictor.

ranges <- sapply(Auto_quant_preds, range)
tibble::tibble(
  Predictor = colnames(ranges),
  Min       = ranges[1, ],
  Max       = ranges[2, ],
  Range     = Max - Min
)

```

```
)
## # A tibble: 6 x 4
##   Predictor      Min      Max   Range
##   <chr>        <dbl>    <dbl>    <dbl>
## 1 cylinders     3       8       5
## 2 displacement  68     455     387
## 3 horsepower    46     230     184
## 4 weight       1613   5140    3527
## 5 acceleration  8      24.8    16.8
## 6 year         70      82      12
rm(ranges)
```

Problem 2 (c)

```
# Compute mean and standard deviation for each.

tibble::tibble(
  Predictor = names(Auto_quant_preds),
  Mean      = sapply(Auto_quant_preds, mean),
  SD        = sapply(Auto_quant_preds, sd)
)

## # A tibble: 6 x 3
##   Predictor      Mean      SD
##   <chr>        <dbl>    <dbl>
## 1 cylinders    5.47    1.71
## 2 displacement 194.    105.
## 3 horsepower   104.    38.5
## 4 weight       2978.   849.
## 5 acceleration 15.5    2.76
## 6 year         76.0    3.68
```

Problem 2 (d)

```
# Remove 10th through 85th observations (inclusive).
Auto_quant_subset <- Auto_quant_preds[-c(10:85), ]

# Compute range, mean, and standard deviation for each predictor on the subset.
ranges_sub <- sapply(Auto_quant_subset, range)
tibble::tibble(
  Predictor = colnames(ranges_sub),
  Min = ranges_sub[1, ],
  Max = ranges_sub[2, ],
  Range = Max - Min,
  Mean = sapply(Auto_quant_subset, mean),
  SD = sapply(Auto_quant_subset, sd)
)

## # A tibble: 6 x 6
##   Predictor     Min     Max   Range     Mean     SD
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 cylinders     3       8       5      5.37    1.65
## 2 displacement  68     455     387    187.     99.7
## 3 horsepower     46     230     184    101.     35.7
## 4 weight      1649    4997    3348   2936.    811.
## 5 acceleration   8.5    24.8    16.3    15.7    2.69
## 6 year        70      82      12     77.1    3.11
rm(Auto_quant_preds, Auto_quant_subset, ranges_sub)
```

Problem 2 (e)

```
# Factor origin.
Auto_plot <- Auto %>%
  dplyr::mutate(origin = factor(origin, labels = c('US', 'EU', 'JP')))

# Inspect pairwise relationships.
```

```
GGally::ggpairs(
  Auto_plot,
  columns = setdiff(colnames(Auto_plot), c('origin', 'name')),
  aes(color = origin, alpha = 0.5)
)
```



```
rm(Auto_plot)
```

```
# Count number of unique names.
```

```
sum(table(Auto$name) == 1)
```

```
## [1] 245
```

The scatterplot matrix shows strong collinearity among the “size/power” predictors: cylinders, displacement, horsepower, and weight move together very tightly (e.g., cylinders-displacement has a correlation around 0.95, and displacement-weight around 0.93), suggesting these variables are largely measuring the same underlying concept (bigger engines/cars tend to be heavier and more powerful). In contrast, acceleration tends to be negatively associated with those size/power variables (most notably

with `horsepower`, around -0.69, and with `displacement`, around -0.54), indicating that cars with larger engines and greater power/weight tend to have smaller `acceleration` values in this dataset. The variable `year` is moderately negatively related to the size/power measures (roughly -0.31 to -0.42 with `weight`, `displacement`, and `horsepower`) and mildly positively related to `acceleration` (about 0.29), consistent with cars becoming lighter and less “big-engine” over time. Finally, the color-group patterns by `origin` suggest systematic differences across regions (U.S. cars clustering at higher weight/displacement/horsepower), and the within-`origin` correlations sometimes differ (e.g., the `cylinders`-`acceleration` relationship is much stronger for U.S. cars than for European or Japanese cars), reinforcing that relationships among predictors can vary by subgroup even when the overall trend is clear. The `name` variable has an extremely large number of unique values (close to the number of observations), so it behaves mostly like an identifier rather than a reusable predictor. Treating it as a categorical feature would create a very high-cardinality factor with many rare levels, which tends to overfit and won’t generalize well.

Problem 2 (f)

Yes. The plots suggest that several variables should be useful for predicting `mpg`. In particular, `mpg` has strong negative associations with `weight`, `displacement`, `horsepower`, and `cylinders` (heavier, larger-engine, higher-power cars tend to have lower gas mileage), so these predictors should have substantial predictive value. The plots also show that `mpg` increases with `year`, indicating that newer model years are generally more fuel-efficient, and `mpg` has a more moderate positive relationship with `acceleration`. Finally, the clear separation by `origin` in the panels suggests that `origin` is also informative: cars from different regions cluster at different typical fuel-efficiency levels even after accounting for other predictors, so it may help explain additional variation in `mpg` beyond the purely quantitative variables.

Problem 3 (a)

```
# Load data.

data(Boston)

# Learn data.

# ?Boston

# Check dimensions.

dim(Boston)
```

```
## [1] 506 13
```

There are 506 rows and 13 columns in the `Boston` data set. Each row corresponds to a census tract/suburb (help file says suburb; textbook says census tract) in the Boston area, and each column is a variable recorded for that tract. One of the columns is `medv`, the median home value (in \$1000s), and the remaining columns are predictors describing characteristics of the tract (e.g., crime rate, number of rooms, property tax rate, etc.).

Problem 3 (b)

```
# Factor origin.

Boston_plot <- Boston %>%
  dplyr::mutate(chas = factor(chas, levels = c(0, 1), labels = c('No River', 'River Bound')))

# Inspect pairwise relationships.

GGally::ggpairs(
  Boston_plot,
  columns = setdiff(colnames(Boston_plot), c('medv', 'chas')),
  aes(color = chas, alpha = 0.5)
)
```



Many predictors are highly skewed, especially `crim`, `zn`, `rad`, and `tax`, with most observations near small values and a handful of extreme outliers, so relationships are often driven by a small number of high-leverage tracts. There is clear collinearity among “urban/industrial” variables: `indus` is strongly positively associated with `nox`, while `dis` is strongly negatively associated with both `indus` and `nox`, suggesting that more industrial and higher-pollution tracts tend to lie closer to employment centers. The variable `age` also tends to be higher where `dis` is lower, consistent with older housing stock being concentrated nearer the city. In addition, `rad` and `tax` are very strongly positively related, indicating that tracts with greater accessibility to radial highways tend to have higher property-tax rates. Finally, `lstat` is positively related to several “urban stress” measures (e.g., `crim`, `nox`, `tax`) and negatively related to variables associated with more desirable suburban tracts (e.g., `zn`, `dis`), reinforcing that multiple predictors are capturing overlapping aspects of neighborhood socioeconomic status and urbanization. The binary predictor `chas` (river adjacency) naturally appears as two bands; any differences involving `chas` are best interpreted as group-level shifts rather

than a continuous linear trend.

Problem 3 (c)

Because `crim` is extremely right-skewed, the scatterplots look compressed and many relationships appear weak visually. Still, the correlation panel indicates moderate associations: `crim` is most positively related to `rad` and `tax` (and also `lstat`, `nox`, and `indus`), and it is moderately negatively related to `dis`. Other predictors (e.g., `zn`, `rm`) show weaker negative associations. Overall, the evidence suggests some predictors are associated with crime rate, but the relationships are not uniformly strong and are influenced by skew/outliers.

Problem 3 (d)

```
# Count top 5 values of variable.

top5_value_counts <- function(df, var) {

  df %>%
    count(Value = .data[[var]], name = 'num_tracts') %>% # frequency table
    arrange(desc(Value)) %>%                                # highest values first
    slice(1:5) %>%                                         # top 5 distinct values
    mutate(Variable = var, .before = 1)
}

top5_value_counts(Boston, 'crim')

##   Variable Value num_tracts
## 1      crim 88.9762        1
## 2      crim 73.5341        1
## 3      crim 67.9208        1
## 4      crim 51.1358        1
## 5      crim 45.7461        1

top5_value_counts(Boston, 'tax')

##   Variable Value num_tracts
## 1      tax    711         5
## 2      tax    666        132
## 3      tax    469         1
## 4      tax    437        15
```

```

## 5      tax    432      9
top5_value_counts(Boston, 'ptratio')

##   Variable Value num_tracts
## 1  ptratio  22.0       2
## 2  ptratio  21.2      15
## 3  ptratio  21.1       1
## 4  ptratio  21.0      27
## 5  ptratio  20.9      11

# Compute ranges.

ranges <- sapply(Boston %>% dplyr::select(crim, tax, ptratio), range)
tibble::tibble(
  Variable = colnames(ranges),
  Min      = ranges[1, ],
  Max      = ranges[2, ],
  Range    = Max - Min
)

## # A tibble: 3 x 4
##   Variable     Min     Max Range
##   <chr>     <dbl>   <dbl> <dbl>
## 1 crim      0.00632  89.0  89.0
## 2 tax        187     711    524
## 3 ptratio    12.6     22     9.4

rm(ranges, top5_value_counts)

```

For `crim`, yes, there are a handful of tracts with extremely high crime rates. The five largest observed `crim` values are approximately 88.98, 73.53, 67.92, 51.14, and 45.75, and each of these occurs only once (so they are isolated extreme outliers). In contrast, the minimum `crim` is about 0.006, so the range is enormous (roughly 89), indicating that crime varies dramatically across tracts and is dominated by a few very high-crime locations. For `tax`, there are also tracts with particularly high values, but the pattern is different: the maximum `tax` rate is 711 and it occurs in 5 tracts, while the next-highest value 666 occurs in a very large number of tracts (132). This indicates that high `tax` rates are not just isolated outliers; there are

sizeable clusters of tracts at very high `tax` levels. The overall range is still large (711 down to 187, range 524), reflecting substantial variation across towns/tracts. For `ptratio`, the highest pupil-teacher ratio is 22.0 (2 tracts), and other top values like 21.2 (15 tracts) and 21.0 (27 tracts) occur for many tracts, so “high” pupil-teacher ratios are fairly common compared to the extreme crime values. However, the spread is much smaller overall: `ptratio` ranges from 12.6 to 22.0 (range 9.4), which is narrow relative to `crim` and `tax`. Overall, `crim` and `tax` show much more dramatic variability (with `crim` having a few extreme outliers and `tax` having large high-value clusters), whereas `ptratio` varies less and its high values are shared by many tracts.

Problem 3 (e)

```
# Count tract(s) bounding Charles River.

sum(Boston$chas == 1)

## [1] 35
```

Problem 3 (f)

```
# Compute median pupil-teacher ratio.

median(Boston$ptratio, na.rm = TRUE)

## [1] 19.05
```

Problem 3 (g)

```
# Identify tract(s) with lowest medv.

Boston_id <- Boston %>%
  dplyr::mutate(tract = dplyr::row_number())
min_medv_tracts <- Boston_id %>% dplyr::slice_min(medv, n = 1, with_ties = TRUE)
min_medv_tracts

##      crim zn indus chas   nox     rm age     dis rad tax ptratio lstat medv tract
## 1 38.3518  0 18.1     0 0.693 5.453 100 1.4896  24 666    20.2 30.59     5  399
## 2 67.9208  0 18.1     0 0.693 5.683 100 1.4254  24 666    20.2 22.98     5  406

# Define helper to make comparison tables.

make_comp_tbl <- function(tid, Boston_id, pred_names, ranges) {
```

```

row <- Boston_id %>% dplyr::filter(tract == tid)
v   <- as.numeric(row[1, pred_names])
tibble::tibble(
  Predictor = pred_names,
  Value      = v,
  Min        = ranges[1, ],
  Max        = ranges[2, ],
  Range      = Max - Min
)
}

# Compare values.

tract_ids <- min_medv_tracts$tract
pred_names <- setdiff(names(Boston), 'medv')
ranges     <- sapply(Boston %>% dplyr::select(dplyr::all_of(pred_names)), range)
make_comp_tbl(399, Boston_id, pred_names, ranges)

## # A tibble: 12 x 5
##   Predictor  Value    Min    Max  Range
##   <chr>     <dbl>  <dbl>  <dbl>  <dbl>
## 1 crim      38.4  0.00632  89.0  89.0
## 2 zn         0     0       100    100
## 3 indus     18.1  0.46    27.7  27.3
## 4 chas      0     0       1      1
## 5 nox       0.693 0.385   0.871  0.486
## 6 rm        5.45  3.56    8.78   5.22
## 7 age       100   2.9     100    97.1
## 8 dis       1.49  1.13    12.1   11.0
## 9 rad       24    1       24     23
## 10 tax      666   187    711    524
## 11 ptratio  20.2  12.6   22     9.4
## 12 lstat    30.6  1.73   38.0   36.2

```

```
make_comp_tbl(406, Boston_id, pred_names, ranges)
```

```
## # A tibble: 12 x 5
##   Predictor Value    Min    Max Range
##   <chr>     <dbl>  <dbl>  <dbl> <dbl>
## 1 crim      67.9  0.00632 89.0  89.0
## 2 zn         0     0       100   100
## 3 indus     18.1  0.46   27.7  27.3
## 4 chas       0     0       1     1
## 5 nox        0.693 0.385  0.871 0.486
## 6 rm         5.68  3.56   8.78  5.22
## 7 age        100   2.9    100   97.1
## 8 dis        1.43  1.13   12.1  11.0
## 9 rad        24    1      24    23
## 10 tax       666   187   711   524
## 11 ptratio   20.2  12.6   22    9.4
## 12 lstat     23.0  1.73   38.0  36.2
```

```
rm(min_medv_tracts, make_comp_tbl, tract_ids, pred_names, ranges)
```

The lowest median home value `medv` occurs in two census tracts, 399 and 406, both with `medv = 5` (in \$1000s). Both tracts share a very similar high-stress urban profile, which helps explain why their home values are so low. In both tracts, the housing stock is as old as it gets in this dataset (`age = 100`, the maximum), they are very close to the city/employment core (`dis = 1.43-1.49`, near the minimum), and they sit in a heavily urban/industrial environment (`indus = 18.1` and `nox = 0.693`, both high). They also have the maximum highway accessibility (`rad = 24`, the dataset maximum) and very high property tax (`tax = 666`, near the top of the range), consistent with the strong `rad-tax` relationship seen elsewhere in the dataset. Where the two tracts differ most is the severity of crime and socioeconomic stress. Tract 399 has `crim = 38.35` together with an extremely high `lstat = 30.59` (near the upper end of the dataset), while tract 406 has even more extreme crime (`crim = 67.92`, among the very highest values) with still-elevated `lstat = 22.98`. In addition, both tracts have relatively low numbers of rooms per dwelling (`rm = 5.45` and `rm = 5.68`), suggesting smaller/less valuable housing compared with higher-`medv` areas.

Problem 3 (h)

```
# Count tracts with rm > 7 and rm > 8.

tibble::tibble(
  Condition = c('rm > 7', 'rm > 8'),
  Count      = c(sum(Boston$rm > 7), sum(Boston$rm > 8))
)

## # A tibble: 2 x 2
##   Condition Count
##   <chr>     <int>
## 1 rm > 7      64
## 2 rm > 8      13

# Inspects tracts with rm > 8.

Boston_id %>%
  dplyr::relocate(tract, .before = 1) %>%
  dplyr::filter(rm > 8) %>%
  dplyr::arrange(dplyr::desc(rm))

##    tract      crim zn indus chas      nox      rm      age      dis rad tax ptratio lstat
## 1    365 3.47428  0 18.10     1 0.7180 8.780 82.9 1.9047  24 666 20.2 5.29
## 2    226 0.52693  0  6.20     0 0.5040 8.725 83.0 2.8944   8 307 17.4 4.63
## 3    258 0.61154 20  3.97     0 0.6470 8.704 86.9 1.8010   5 264 13.0 5.12
## 4    263 0.52014 20  3.97     0 0.6470 8.398 91.5 2.2885   5 264 13.0 5.91
## 5    164 1.51902  0 19.58     1 0.6050 8.375 93.9 2.1620   5 403 14.7 3.32
## 6    233 0.57529  0  6.20     0 0.5070 8.337 73.3 3.8384   8 307 17.4 2.47
## 7    268 0.57834 20  3.97     0 0.5750 8.297 67.0 2.4216   5 264 13.0 7.44
## 8    225 0.31533  0  6.20     0 0.5040 8.266 78.3 2.8944   8 307 17.4 4.14
## 9    254 0.36894 22  5.86     0 0.4310 8.259  8.4 8.9067   7 330 19.1 3.54
## 10   234 0.33147  0  6.20     0 0.5070 8.247 70.4 3.6519   8 307 17.4 3.95
## 11   98  0.12083  0  2.89     0 0.4450 8.069 76.0 3.4952   2 276 18.0 4.21
## 12   227 0.38214  0  6.20     0 0.5040 8.040 86.5 3.2157   8 307 17.4 3.13
## 13   205 0.02009 95  2.68     0 0.4161 8.034 31.9 5.1180   4 224 14.7 2.88
##    medv
```

```

## 1 21.9
## 2 50.0
## 3 50.0
## 4 48.8
## 5 50.0
## 6 41.7
## 7 50.0
## 8 44.8
## 9 42.8
## 10 48.3
## 11 38.7
## 12 37.6
## 13 50.0

```

```
rm(Boston_id)
```

In this data set, 64 census tracts have an average of more than 7 rooms per dwelling (`rm > 7`), and 13 census tracts have an average of more than 8 rooms per dwelling (`rm > 8`). The `rm > 8` tracts are relatively rare and, based on the printed rows and contrasting with the roughly estimated ranges deduced from the pair plots in part (b), generally resemble higher-end neighborhoods: many have very high median home values (`medv` is frequently near the top of the scale, including several tracts with `medv = 50`), and their `lstat` values tend to be low (roughly 2–7), which is consistent with more affluent areas. Crime rates for these tracts are typically low to moderate (most `crim` values are well below the extreme outliers seen elsewhere in the dataset), though there is at least one clear exception: tract 365 has the largest `rm` value (`rm = 8.78`) but also relatively high `crim` and very high `nox`, showing that very large homes can coexist with some urban disamenities. Overall, tracts with `rm > 8` tend to correspond to especially desirable areas with high home values and favorable socioeconomic indicators, with a small number of exceptions.