

# STA 336: STATISTICAL MACHINE LEARNING

## Homework 4

Rento Saijo

February 24, 2026

### Table of Contents

<a href="#">Disclosure</a>	2
<a href="#">Problem 1</a>	3
<a href="#">Problem 2</a>	4
<a href="#">Problem 3</a>	5
<a href="#">Problem 4</a>	7
<a href="#">Problem 5</a>	8
<a href="#">Problem 6</a>	9
<a href="#">Problem 7</a>	11
<a href="#">Problem 8</a>	13
<a href="#">Problem 9</a>	14
<a href="#">Problem 10</a>	15
<a href="#">Problem 11</a>	17
<a href="#">Problem 12</a>	18
<a href="#">Problem 13</a>	19
<a href="#">Problem 14</a>	20

## Disclosure

GPT-5.3-Codex was used to create the YAML portion and some `LaTeX` code to format the text/equations nicely. Page formatting code was also provided by Derin Gezgin. In the setup chunk, libraries were loaded and some helper functions were defined including but not limited to `table_latex()` and `with_family()`. See the original RMD file [here](#) for more details.

## Problem 1

### PROBLEM 1

Use the full dataset and build a model using a single predictor on student. What is the mathematical form of the model?

```
Cell 1
1 # Load data.
2 default_df <- ISLR2::Default
3
4 # Fit student-only logistic model.
5 model_q1 <- stats::glm(default ~ student, data = default_df, family =
6   stats::binomial)
7
8 # Build coefficient table.
9 coef_q1_tbl <- tibble::tibble(
10   term      = names(stats::coef(model_q1)),
11   estimate  = as.numeric(stats::coef(model_q1))
12 ) %>%
13   dplyr::mutate(estimate = round(estimate, 6))
14
15 # Print coefficient table.
16 coef_q1_tbl %>%
17   table_latex(
18     col_names = c('Term', 'Estimate'),
19     caption  = 'Logistic model coefficients with student only.'
)
```

Table 1: Logistic model coefficients with student only.

Term	Estimate
(Intercept)	-3.504128
studentYes	0.404887

The fitted model is

$$\log\left(\frac{\hat{p}}{1 - \hat{p}}\right) = -3.504128 + 0.404887 \cdot I\{\text{student} = \text{Yes}\},$$

where  $\hat{p} = \Pr(\text{default} = \text{Yes} | \text{student})$ . Equivalently,

$$\hat{p} = \frac{\exp(-3.504128 + 0.404887 \cdot I\{\text{student} = \text{Yes}\})}{1 + \exp(-3.504128 + 0.404887 \cdot I\{\text{student} = \text{Yes}\})}.$$

## Problem 2

### PROBLEM 2

How would you interpret the coefficient of student on the model of Problem 1?

The coefficient on `studentYes` is 0.404887, so the log-odds of default are higher for students than non-students in this one-predictor model. The corresponding odds ratio is

$$\exp(0.404887) = 1.4991.$$

Thus, the model estimates that students have about  $1.50\times$  the odds of default relative to non-students.

## Problem 3

### PROBLEM 3

What is the confusion matrix when you use the cutoff point to be 0.5? What is the error rate of prediction?

Cell 2

```

20 # Build predictions with 0.5 cutoff.
21 prob_q1 <- stats::predict(model_q1, type = 'response')
22 pred_q1 <- factor(ifelse(prob_q1 > 0.5, 'Yes', 'No'), levels = c('No', 'Yes'))
23 actual_q1 <- factor(default_df$default, levels = c('No', 'Yes'))
24
25 # Build confusion matrix.
26 cm_q1 <- table(actual = actual_q1, predicted = pred_q1)
27 cm_q1_tbl <- as.data.frame.matrix(cm_q1) %>%
28   tibble::rownames_to_column(var = 'Actual')
29
30 # Compute error.
31 error_q1 <- mean(pred_q1 != actual_q1)
32 error_q1_tbl <- tibble::tibble(
33   metric = 'Prediction error rate',
34   value = round(error_q1, 4)
35 )
36
37 # Print tables.
38 cm_q1_tbl %>%
39   table_latex(
40   col_names = c('Actual', 'Predicted No', 'Predicted Yes'),
41   caption = 'Confusion matrix for student-only model (cutoff = 0.5).'
42 )

```

Table 2: Confusion matrix for student-only model (cutoff = 0.5).

Actual	Predicted No	Predicted Yes
No	9667	0
Yes	333	0

Cell 3

```

43 error_q1_tbl %>%
44   table_latex(
45   col_names = c('Metric', 'Value'),
46   caption = 'Error rate for student-only model (cutoff = 0.5).'
47 )

```

Table 3: Error rate for student-only model (cutoff = 0.5).

Metric	Value
Prediction error rate	0.0333

At cutoff 0.5, this model predicts every observation as No. The prediction error rate is 0.0333 (about 3.33%).

## Problem 4

### PROBLEM 4

Use the full dataset and build a model using a single predictor on balance. What is the mathematical form of the model?

```
Cell 4
48 # Fit balance-only logistic model.
49 model_q4 <- stats::glm(default ~ balance, data = default_df, family =
50   stats::binomial)
51
52 # Build coefficient table.
53 coef_q4_tbl <- tibble::tibble(
54   term      = names(stats::coef(model_q4)),
55   estimate  = as.numeric(stats::coef(model_q4))
56 ) %>%
57   dplyr::mutate(estimate = round(estimate, 6))
58
59 # Print coefficient table.
60 coef_q4_tbl %>%
61   table_latex(
62     col_names = c('Term', 'Estimate'),
63     caption  = 'Logistic model coefficients with balance only.'
64   )
```

Table 4: Logistic model coefficients with balance only.

Term	Estimate
(Intercept)	-10.651331
balance	0.005499

The fitted model is

$$\log\left(\frac{\hat{p}}{1 - \hat{p}}\right) = -10.651331 + 0.005499 \text{balance},$$

where  $\hat{p} = \Pr(\text{default} = \text{Yes} \mid \text{balance})$ . Equivalently,

$$\hat{p} = \frac{\exp(-10.651331 + 0.005499 \text{balance})}{1 + \exp(-10.651331 + 0.005499 \text{balance})}.$$

## Problem 5

### PROBLEM 5

How would you interpret the coefficient of balance on the model of Problem 4?

The coefficient on **balance** is 0.005499, so each one-unit increase in balance increases the log-odds of default by 0.005499. In odds terms,

$$\exp(0.005499) = 1.0055,$$

so each \$1 increase in balance multiplies the odds of default by about 1.0055.

## Problem 6

### PROBLEM 6

What is the confusion matrix of the model in Problem 4 when you use the cutoff point to be 0.5? What is the error rate of prediction?

Cell 5

```

64 # Build predictions with 0.5 cutoff.
65 prob_q4 <- stats::predict(model_q4, type = 'response')
66 pred_q4 <- factor(ifelse(prob_q4 > 0.5, 'Yes', 'No'), levels = c('No', 'Yes'))
67 actual_q4 <- factor(default_df$default, levels = c('No', 'Yes'))
68
69 # Build confusion matrix.
70 cm_q4 <- table(actual = actual_q4, predicted = pred_q4)
71 cm_q4_tbl <- as.data.frame.matrix(cm_q4) %>%
72   tibble::rownames_to_column(var = 'Actual')
73
74 # Compute error.
75 error_q4 <- mean(pred_q4 != actual_q4)
76 error_q4_tbl <- tibble::tibble(
77   metric = 'Prediction error rate',
78   value = round(error_q4, 4)
79 )
80
81 # Print tables.
82 cm_q4_tbl %>%
83   table_latex(
84     col_names = c('Actual', 'Predicted No', 'Predicted Yes'),
85     caption = 'Confusion matrix for balance-only model (cutoff = 0.5).'
86   )

```

Table 5: Confusion matrix for balance-only model (cutoff = 0.5).

Actual	Predicted No	Predicted Yes
No	9625	42
Yes	233	100

Cell 6

```

87 error_q4_tbl %>%
88   table_latex(
89   col_names = c('Metric', 'Value'),
90   caption = 'Error rate for balance-only model (cutoff = 0.5).'
91 )

```

Table 6: Error rate for balance-only model (cutoff = 0.5).

Metric	Value
Prediction error rate	0.0275

The balance-only model has prediction error 0.0275 (about 2.75%).

## Problem 7

### PROBLEM 7

Split the data into halves and use one half as the training and the other half as the test. Fit a model on the training data using `student`, `balance`, and `income` as predictors.

Cell 7

```

92 # Create train/test split.
93 set.seed(20060527)
94 n_default      <- nrow(default_df)
95 train_idx       <- sample(seq_len(n_default), n_default / 2)
96 default_train  <- default_df[train_idx, ]
97 default_test   <- default_df[-train_idx, ]
98
99 # Fit training model.
100 model_q7 <- stats::glm(
101   default ~ student + balance + income,
102   data    = default_train,
103   family  = stats::binomial
104 )
105
106 # Build split table.
107 split_q7_tbl <- tibble::tibble(
108   data_set     = c('Training', 'Test'),
109   observations = c(nrow(default_train), nrow(default_test))
110 )
111
112 # Build coefficient table.
113 coef_q7_tbl <- as.data.frame(stats::coef(summary(model_q7))) %>%
114   tibble::rownames_to_column(var = 'term') %>%
115   dplyr::rename(
116     estimate  = Estimate,
117     std_error = `Std. Error`,
118     z_value   = `z value`,
119     p_value   = `Pr(>|z|)`
120   ) %>%
121   dplyr::mutate(dplyr::across(-term, ~ round(.x, 6)))
122
123 # Create probabilities for later questions.
124 prob_train_q7 <- stats::predict(model_q7, newdata = default_train, type =
125   'response')
126 prob_test_q7  <- stats::predict(model_q7, newdata = default_test, type =
127   'response')
128
129 # Print tables.
130 split_q7_tbl %>%
131   table_latex(
132     col_names = c('Data set', 'Observations'),
133     caption  = 'Train/test split sizes.'
```

132 )

Table 7: Train/test split sizes.

Data set	Observations
Training	5000
Test	5000

---

Cell 8

---

```

133 coef_q7_tbl %>%
134   table_latex(
135     col_names = c('Term', 'Estimate', 'Std. Error', 'z value', 'Pr(>|z|)'),
136     caption  = 'Logistic regression coefficient summary.'
137   )

```

Table 8: Logistic regression coefficient summary.

Term	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-11.362494	0.729852	-15.568209	0.000000
studentYes	-0.439823	0.345068	-1.274598	0.202452
balance	0.005885	0.000340	17.319003	0.000000
income	0.000008	0.000012	0.693137	0.488224

## Problem 8

### PROBLEM 8

What is the training error of the model in Problem 7?

```
Cell 9
138 # Build training predictions.
139 pred_train_q7 <- factor(ifelse(prob_train_q7 > 0.5, 'Yes', 'No'), levels =
140   c('No', 'Yes'))
141 actual_train_q7 <- factor(default_train$default, levels = c('No', 'Yes'))
142
143 # Compute training error.
144 train_error_q7 <- mean(pred_train_q7 != actual_train_q7)
145 train_error_tbl <- tibble::tibble(
146   metric = 'Training error rate',
147   value = round(train_error_q7, 4)
148 )
149
150 # Print training error table.
151 train_error_tbl %>%
152   table_latex(
153     col_names = c('Metric', 'Value'),
154     caption = 'Training error with cutoff = 0.5.'
155   )
```

Table 9: Training error with cutoff = 0.5.

Metric	Value
Training error rate	0.0264

The training error is 0.0264 (about 2.64%).

## Problem 9

### PROBLEM 9

What is the test error of the model in Problem 7?

Cell 10

```

155 # Build test predictions.
156 pred_test_q7 <- factor(ifelse(prob_test_q7 > 0.5, 'Yes', 'No'), levels =
157   c('No', 'Yes'))
158 actual_test_q7 <- factor(default_test$default, levels = c('No', 'Yes'))
159
160 # Compute test error.
161 test_error_q7 <- mean(pred_test_q7 != actual_test_q7)
162 test_error_tbl <- tibble::tibble(
163   metric = 'Test error rate',
164   value = round(test_error_q7, 4)
165 )
166
167 # Print test error table.
168 test_error_tbl %>%
169   table_latex(
170     col_names = c('Metric', 'Value'),
171     caption = 'Test error with cutoff = 0.5.'
172   )

```

Table 10: Test error with cutoff = 0.5.

Metric	Value
Test error rate	0.0266

The test error is 0.0266 (about 2.66%).

## Problem 10

### PROBLEM 10

Draw the ROC curve for model in Problem 7 and find the best cutoff point. What does the best cutoff point mean?

Cell 11

```

172 # Build ROC.
173 roc_q7 <- Epi::ROC(
174   form = default ~ student + balance + income,
175   data = default_train,
176   plot = 'ROC',
177   MX   = TRUE
178 )

```

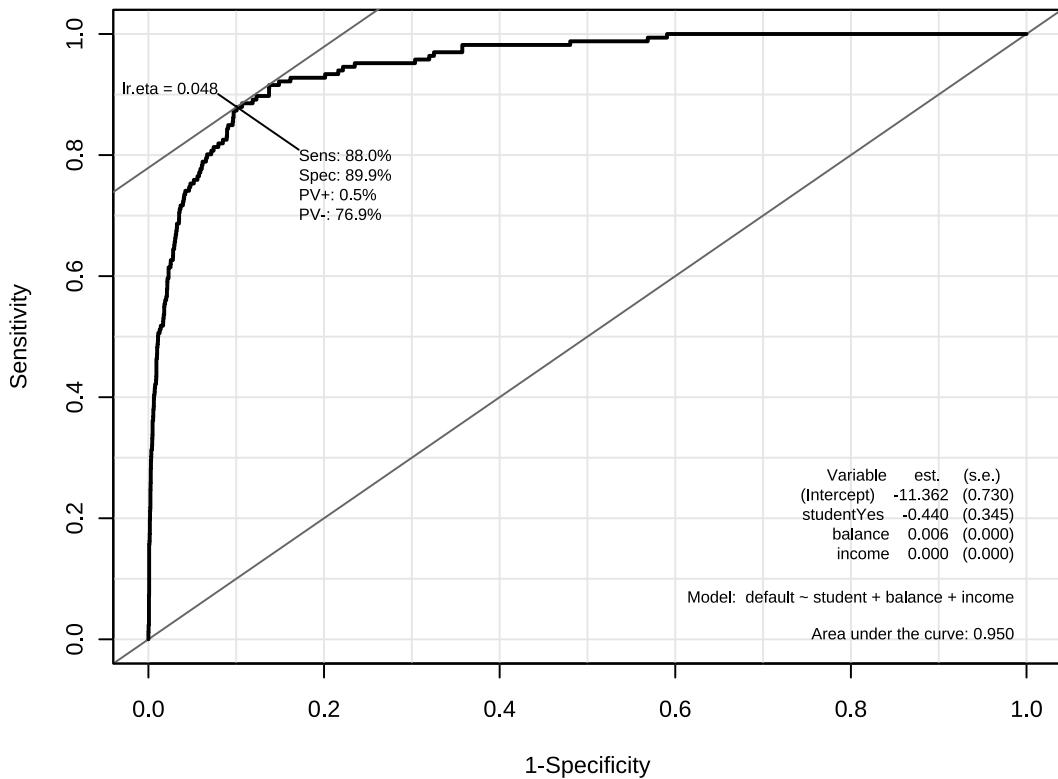


Figure 1: ROC curve on training data for the multivariable model.

Cell 12

```

179 # Extract ROC grid.
180 roc_q7_res <- roc_q7$res
181
182 # Find best cutoff with MX criterion.

```

```

183 best_row_q7 <- roc_q7_res %>%
184   dplyr::mutate(mx_value = sens + spec) %>%
185   dplyr::filter(mx_value == max(mx_value)) %>%
186   dplyr::arrange(dplyr::desc(lr.eta)) %>%
187   dplyr::slice(1)
188 best_cutoff_q7 <- best_row_q7$lr.eta[[1]]
189
190 # Build best-cutoff table.
191 best_cutoff_tbl <- tibble::tibble(
192   metric = c(
193     'Best cutoff',
194     'Sensitivity',
195     'Specificity',
196     'False positive rate',
197     'MX value (sensitivity + specificity)',
198     'AUC'
199   ),
200   value = round(c(
201     best_cutoff_q7,
202     best_row_q7$sens[[1]],
203     best_row_q7$spec[[1]],
204     1 - best_row_q7$spec[[1]],
205     best_row_q7$mx_value[[1]],
206     roc_q7$AUC
207   ), 4)
208 )
209
210 # Print best-cutoff table.
211 best_cutoff_tbl %>%
212   table_latex(
213     col_names = c('Metric', 'Value'),
214     caption = 'Best cutoff chosen from training ROC.'
215   )

```

Table 11: Best cutoff chosen from training ROC.

Metric	Value
Best cutoff	0.0477
Sensitivity	0.8795
Specificity	0.8995
False positive rate	0.1005
MX value (sensitivity + specificity)	1.7790
AUC	0.9498

Using the class ROC convention (`Epi::ROC` with `MX = TRUE`), the best cutoff is 0.0477. This is the threshold that maximizes sensitivity + specificity on the training data.

## Problem 11

### PROBLEM 11

Use the best cutoff point on the test data. What is the test error now?

Cell 13

```

216 # Build test predictions with best cutoff.
217 pred_test_best_q7 <- factor(ifelse(prob_test_q7 > best_cutoff_q7, 'Yes', 'No'),
218   levels = c('No', 'Yes'))
219
220 # Compute test error with best cutoff.
221 test_error_best_q7 <- mean(pred_test_best_q7 != actual_test_q7)
222 test_error_best_tbl <- tibble::tibble(
223   metric = 'Test error rate (best cutoff)',
224   value = round(test_error_best_q7, 4)
225 )
226
227 # Print test error table.
228 test_error_best_tbl %>%
229   table_latex(
230     col_names = c('Metric', 'Value'),
231     caption = 'Test error using best cutoff from training ROC.'
232   )

```

Table 12: Test error using best cutoff from training ROC.

Metric	Value
Test error rate (best cutoff)	0.1016

Using the best cutoff from Problem 10, the test error is 0.1016 (about 10.16%).

## Problem 12

### PROBLEM 12

For a student, with 0 balance, and 20000 income, what is the predicted probability of default? Use the model on training data.

Cell 14

```

232 # Build student profile.
233 new_student_q12 <- data.frame(
234   student = factor('Yes', levels = levels(default_train$student)),
235   balance = 0,
236   income   = 20000
237 )
238
239 # Compute predicted probability.
240 p_q12 <- as.numeric(stats::predict(model_q7, newdata = new_student_q12, type =
241   'response'))
242 q12_tbl <- tibble::tibble(
243   metric = 'Predicted probability',
244   value   = signif(p_q12, 6)
245 )
246
247 # Print probability table.
248 q12_tbl %>%
249   table_latex(
250     col_names = c('Metric', 'Value'),
251     caption  = 'Predicted probability for student profile.'
252   )

```

Table 13: Predicted probability for student profile.

Metric	Value
Predicted probability	0.0000088

The predicted probability is 0.00000882 (approximately).

## Problem 13

### PROBLEM 13

For a non-student, with 0 balance, and 20000 income, what is the predicted probability of default? Use the model on training data.

Cell 15

```

252 # Build non-student profile.
253 new_nonstudent_q13 <- data.frame(
254   student = factor('No', levels = levels(default_train$student)),
255   balance = 0,
256   income   = 20000
257 )
258
259 # Compute predicted probability.
260 p_q13 <- as.numeric(stats::predict(model_q7, newdata = new_nonstudent_q13, type =
261   'response'))
262 q13_tbl <- tibble::tibble(
263   metric = 'Predicted probability',
264   value   = signif(p_q13, 6)
265 )
266
267 # Print probability table.
268 q13_tbl %>%
269   table_latex(
270     col_names = c('Metric', 'Value'),
271     caption  = 'Predicted probability for non-student profile.'
272   )

```

Table 14: Predicted probability for non-student profile.

Metric	Value
Predicted probability	0.0000137

The predicted probability is 0.00001369 (approximately).

## Problem 14

### PROBLEM 14

What is the difference between the predicted values of Problems 12 and 13? Does this difference come from the coefficient of the student variable? How are they related?

Cell 16

```

272 # Compute difference and student effect.
273 diff_q14 <- p_q12 - p_q13
274 beta_student_q14 <- as.numeric(stats::coef(model_q7)['studentYes'])
275 odds_ratio_student_q14 <- exp(beta_student_q14)
276
277 # Build relation table.
278 q14_tbl <- tibble::tibble(
279   metric = c(
280     'p(student = Yes) - p(student = No)',
281     'student coefficient (log-odds scale)',
282     'student odds ratio'
283   ),
284   value = c(diff_q14, beta_student_q14, odds_ratio_student_q14)
285 ) %>%
286   dplyr::mutate(value = signif(value, 6))
287
288 # Print relation table.
289 q14_tbl %>%
290   table_latex(
291     col_names = c('Metric', 'Value'),
292     caption = 'Probability difference and student coefficient relation.'
293   )

```

Table 15: Probability difference and student coefficient relation.

Metric	Value
p(student = Yes) - p(student = No)	-0.0000049
student coefficient (log-odds scale)	-0.4398230
student odds ratio	0.6441500

The difference is  $p_{12} - p_{13} = -4.87 \times 10^{-6}$ , so the student profile has a slightly lower predicted default probability here. Yes, this difference comes from the **studentYes** coefficient. For fixed **balance** and **income**, changing student status shifts the log-odds by  $\hat{\beta}_{\text{studentYes}}$ :

$$\log\left(\frac{\hat{p}_{\text{Yes}}}{1 - \hat{p}_{\text{Yes}}}\right) - \log\left(\frac{\hat{p}_{\text{No}}}{1 - \hat{p}_{\text{No}}}\right) = \hat{\beta}_{\text{studentYes}}.$$

Thus, the odds are multiplied by  $\exp(\hat{\beta}_{\text{studentYes}}) = 0.6442$  when moving from non-student to student at the same **balance** and **income**.